

To view the counties with the highest ambient ozone pollution, we can check the head of the sorted.counties list.

```
head(sorted.counties, 10)
```

```
##      Group.1      Group.2      x
## 148 Colorado Clear Creek 0.05206572
## 10  Wyoming  Albany 0.05101239
## 263 Colorado  Gilpin 0.04829935
## 261 Arizona    Gila 0.04784383
## 607  Utah      San Juan 0.04716827
## 751 Arizona  Yavapai 0.04668674
## 281 Colorado  Gunnison 0.04667247
## 736  Nevada  White Pine 0.04654307
## 336 Colorado  Jefferson 0.04649084
## 156 Arizona   Cochise 0.04591013
```

This would give us the counties with the highest ozone pollution. Compared to the 2020 analysis -which I checked from the lecture slides- we can see similar counties (Colorado-Clear Creek, Wyoming-Albany, Colorado-Gilpin) in top-5.

To determine if the hourly ozone levels on average are higher in New York than in California, we can group our data by state and apply the summary function.

```
group.states.summary = aggregate(ozone.2024$Sample.Measurement,
                                by = list(ozone.2024$State.Name),
                                FUN=summary)
```

Following this, we can extract and combine New York and California from this to be able to compare them

```
NY.summary = group.states.summary[group.states.summary$Group.1 == "New
York",]
CA.summary = group.states.summary[group.states.summary$Group.1 ==
"California",]
NY.CA.summary = rbind(NY.summary, CA.summary)
NY.CA.summary

##      Group.1      x.Min.      x.1st Qu.      x.Median      x.Mean      x.3rd Qu.
## 34  New York  0.00000000  0.02300000  0.03200000  0.03103199  0.03900000
## 5   California -0.00400000  0.02200000  0.03300000  0.03299081  0.04300000
##      x.Max.
## 34  0.11500000
## 5   0.14700000
```

When we look at the summary output for New York and California, we can see that California has slightly higher mean hourly ozone levels compared to New York.

To determine if the hourly ozone levels on average are higher in New York City than are in LA using the 5-number summary, we can group the data-frame by counties using the summary function.

```
group.counties.summary = aggregate(ozone.2024$Sample.Measurement,
                                   by = list(ozone.2024$State.Name,
                                             ozone.2024$County.Name),
                                   FUN = summary)
```

Following this, we can simply extract and combine the summaries for New York City and Los Angeles.

```
NYC.summary = group.counties.summary[group.counties.summary$Group.2 == "New
York", ]
LA.summary = group.counties.summary[group.counties.summary$Group.2 == "Los
Angeles", ]
NYC.LA.summary = rbind(NYC.summary, LA.summary)
NYC.LA.summary
```

	Group.1	Group.2	x.Min.	x.1st Qu.	x.Median	x.Mean
## 501	New York	New York	0.00000000	0.01900000	0.02800000	0.02886334
## 413	California	Los Angeles	-0.00400000	0.02000000	0.03300000	0.03326734
	x.3rd Qu.	x.Max.				
## 501	0.03700000	0.10800000				
## 413	0.04400000	0.14300000				

From this table, we can see that Los Angeles has a slightly higher average hourly ozone level compared to the New York City. 20/20

Problem #2

Collecting data is often a messy process resulting in multiple errors in the data. Consider the following small vector representing the weights of 10 adults in pounds.

```
my.weights <- c(150, 138, 289, 239, 12, 103, 310, 200, 218, 178)
```

As far as we know, it's not possible for an adult to weigh 12 pounds, so that is most likely an error. Change this value to NA, and then find the standard deviation of the weights after removing the NA value.

To replace the value 12 with NA we can use the *exactly equal to* operator.

```
my.weights[my.weights == 12] = NA
my.weights
```

## [1]	150	138	289	239	NA	103	310	200	218	178
--------	-----	-----	-----	-----	----	-----	-----	-----	-----	-----

Before calculating the standard deviation, we should remove the NA value from the vector

```
my.weights = my.weights[!is.na(my.weights)]
my.weights
```

```
## [1] 150 138 289 239 103 310 200 218 178
```

Finally, we can use the sd function to calculate the standard deviation

```
sd(my.weights)
```

```
## [1] 68.95067 5/5
```

Problem #3

Consider the following variables: age and income:

```
age <- c("middle age", "senior", "middle age", "senior", "senior", "senior",
"senior", "middle age")
income <- c("lower", "lower", "upper", "middle", "upper", "lower", "lower",
"middle")
```

What is the class of each variable?

To find out what class is age and income, we can use the class function.

```
class(age)
```

```
## [1] "character"
```

```
class(income)
```

```
## [1] "character"
```

We can see that both age and income are in the character class.

Change the age variable to a factor with levels for age as: youth, young adult, middle age, and senior.

To change the age variable in a factor, we can simply use the factor function

```
age.factor = factor(age,
                     levels = c("youth", "young adult", "middle age",
"senior"))
age.factor

## [1] middle age senior      middle age senior      senior      senior
senior
## [8] middle age
## Levels: youth young adult middle age senior
```

Change the income variable to a factor with levels as lower, middle, and upper.

Same as the previous one, to change the income variable to a factor, we can use the factor function

```
income.factor = factor(income,
                        levels = c("lower", "middle", "upper"))
income.factor
## [1] lower lower upper middle upper lower lower middle
## Levels: lower middle upper 10/10
```

Problem #4

Suppose you keep track of the mileage each time you fill up. At your last 8 fill-ups the mileage was: 65311, 65624, 65908, 66219, 66499, 66821, 67145, 67447.

Enter this data into R. Use the function diff on the data (use ?diff). Use the documentation for diff function to learn about it and then explain briefly what does this function give you?

To enter this mileage data into R, we can basically save these values in a vector

```
mileage = c(65311, 65624, 65908, 66219, 66499, 66821, 67145, 67447)
```

When I checked the hep documentation and examples, I saw that the diff function basically takes the difference between each consecutive element in a vector. For example, in this case the elements of diff(mileage) would be

Index	Difference (Next - Current)
1	65624 - 65311 = 313
2	65908 - 65624 = 284
3	66219 - 65908 = 311
4	66499 - 66219 = 280
5	66821 - 66499 = 322
6	67145 - 66821 = 324
7	67447 - 67145 = 302

From the R output, we can see these values as well.

```
diff(mileage)
## [1] 313 284 311 280 322 324 302 5/5
```

Problem #5

Create the following data frame:

```
my_data <- data.frame(student_id = c(100234, 132454, 453123),
                      test_1_grade = c(82, 93, 87),
                      hw_1_grade = c(92, 89, 98),
                      session = c("7 AM", "7 PM", "7 AM"))
```

my_data

```
##   student_id test_1_grade hw_1_grade session
## 1    100234         82         92    7 AM
## 2    132454         93         89    7 PM
## 3    453123         87         98    7 AM
```

Obtain the column names of our data frame.

To get the column names, we can basically use the `colnames` function.

```
colnames(my_data)
```

```
## [1] "student_id" "test_1_grade" "hw_1_grade" "session"
```

This function returns the column names which are `student_id`, `test_1_grade`, `hw_1_grade`, and `session`.

Get the number of rows or columns in a data frame, try `nrow()`, `ncol()`, or `dim()` functions.

```
nrow(my_data)
```

```
## [1] 3
```

```
ncol(my_data)
```

```
## [1] 4
```

```
dim(my_data)
```

```
## [1] 3 4
```

`nrow()` returns the row count of our data-frame

`ncol()` returns the column count of our data-frame

`dim()` returns the dimensions (row count x column count) of our data-frame.

To subset rows and columns of a data frame we can use the following syntax:

my_data_frame[row condition, column condition]. The row/column conditions may be either numeric indexes, logical expressions, or vectors

Explain the subset you get from the following code:

```
my_data_frame <- data.frame(make = c("Toyota", "Honda", "Ford", "Toyota",  
"Ford", "Honda"),  
                             mpg = c(34, 33, 22, 32, 29, 27),  
                             cylinders = c(4, 4, 8, 6, 6, 8))  
  
my_data_frame  
##      make mpg cylinders  
## 1 Toyota  34         4  
## 2  Honda  33         4  
## 3   Ford  22         8  
## 4 Toyota  32         6  
## 5   Ford  29         6  
## 6  Honda  27         8
```

This code snippet creates a new data-frame with 3 columns (make, mpg, cylinders) and 6 rows. This data-frame gives information about the brand, distance the car can go on one gallon of fuel, and the number of cylinders the car has for 6 different cars.

Explain the subset you get from the following code:

```
my_data_frame[1:3, 1:2]  
  
##      make mpg  
## 1 Toyota  34  
## 2  Honda  33  
## 3   Ford  22
```

This code snippet selects rows 1 to 3 and first two columns.

```
my_data_frame[c(1, 2, 3), c(1, 2)]  
  
##      make mpg  
## 1 Toyota  34  
## 2  Honda  33  
## 3   Ford  22
```

Similar to the previous example, this code snippet also selects the first 3 rows and first 2 columns. 1:3 is equal to creating a vector from 1 to 3 which the second code snippet does.

```
my_data_frame[2, c(1, 3)]
```

```
##      make cylinders
## 2 Honda           4
```

This code snippet specifically selects the second row of the data-frame. At the same time, it selects the first and the third columns. Different than the previous examples this is **not** a range in form of 1:3, it is a vector with 1 **and** 3.

```
my_data_frame[1:2, ]

##      make mpg cylinders
## 1 Toyota  34           4
## 2  Honda  33           4
```

This example takes the first two rows of the data-frame. At the same time, it selects all the columns as nothing was specified in the column selection part.

```
logical_condition <- my_data_frame$mpg >= 30
my_data_frame[logical_condition, ]

##      make mpg cylinders
## 1 Toyota  34           4
## 2  Honda  33           4
## 4 Toyota  32           6
```

This code snippet first defines a logical condition which would return TRUE or FALSE for each row if that the mpg value of that row is more than or equal to 30. The logical_condition variable is a vector with TRUE and FALSE values.

In the second part of the code, this vector is used to filter the rows which has an mpg value less than 30. Similar to the previous example, as the column is not specified, it returns all 3 columns.

```
my_data_frame[my_data_frame$mpg >= 30, ]

##      make mpg cylinders
## 1 Toyota  34           4
## 2  Honda  33           4
## 4 Toyota  32           6
```

This example is exactly identical as the previous example as it again checks if the mpg value for each row is more then or equal to 30. Similarly, it selects all the columns as it does not specify anything related to columns.

```
my_data_frame[my_data_frame$mpg >= 32, c(2, 3)]
```

##	mpg	cylinders
## 1	34	4
## 2	33	4
## 4	32	6

This example merges the previous examples together. We select the rows with an mpg value larger than or equal to 32. At the same time, it selects the 2nd and the 3rd columns by specifying them in a vector.

```
my_data_frame[my_data_frame$mpg >= 32, c("mpg", "cylinders")]
```

##	mpg	cylinders
## 1	34	4
## 2	33	4
## 4	32	6

In this example, it first selects the rows with an mpg value larger then or equal to 32. At the same time, it specifically selects the columns named mpg and cylinders.

Now try to subset `my_data_frame` to only include rows that have a `cylinders` value of 4.

```
my_data_frame[my_data_frame$cylinders == 4, ]
##      make mpg cylinders
## 1 Toyota  34          4
## 2  Honda  33          4
```

To do this, we can basically use the *equal to* operator. We have to select the rows with cylinder values equal to 4. At the same time, we should ensure that we select all the columns.

From the result, we can see that the first and second vehicles (branded Toyota and Honda with 34 and 33 mpg respectively) has 4 cylinders. 15/15

Problem #6

Create data as the following list:

```
my_list <- list(classes_offered = c("MIS 431", "MIS 310", "MIS 410", "MIS  
412"),  
                student_data = data.frame(student_id = c(54, 100, 32, 423, 2,  
19, 39),  
                                           age = c(18, 22, 27, 18, 29, 22,  
20)),
```

```
3.4, 3.2),  
gpa = c(3.1, 2.8, 3.7, 3.4, 3.2,  
stringsAsFactors = FALSE))
```

Write the R code that calculates the median value (use the median() function) of the gpa variable in student_data. All you need to do is pass the student_id vector into the median() function.

This code snippet creates a list of a vector and a data-frame. We should first access to the student_data data-frame which is inside my_list. Following this, we should access to the GPA column in order to pass it to the median function and calculate the median GPA. We can achieve this chain access using the \$ operator.

```
gpa.values = my_list$student_data$gpa # Saving the GPA values  
median(gpa.values) # Calculating the median  
## [1] 3.2 10/10
```

Problem #7

Let us first create dataframes.

```
Feature1A <- c("A", "B", "C", "D")  
Feature2A <- c(1000, 2000, 3000, 4000)  
Feature3A <- c(25.5, 35.5, 45.5, 55.5)  
Feature4A <- c(10, 34, 78, 3)  
Dataframe1 <- data.frame(Feature1A, Feature2A, Feature3A, Feature4A)  
colnames(Dataframe1) <- c("Feature1", "Feature2", "Feature3", "Feature4")  
Dataframe1  
  
##   Feature1 Feature2 Feature3 Feature4  
## 1      A    1000    25.5      10  
## 2      B    2000    35.5      34  
## 3      C    3000    45.5      78  
## 4      D    4000    55.5       3  
  
# creating Dataframe2  
Feature1B <- c("E", "F", "G", "H")  
Feature2B <- c(5000, 6000, 7000, 8000)  
Feature3B <- c(65.5, 75.5, 85.5, 95.5)  
Dataframe2 <- data.frame(Feature1B, Feature2B, Feature3B)  
colnames(Dataframe2) <- c("Feature1", "Feature2", "Feature3")  
Dataframe2  
  
##   Feature1 Feature2 Feature3  
## 1      E    5000    65.5  
## 2      F    6000    75.5
```

```
## 3      G      7000      85.5
## 4      H      8000      95.5
```

Merge merges Features 1-3 of the two data frames and called the resulting dataframe as Output. Use function merge().

```
output = merge(Dataframe1, Dataframe2, all = TRUE)
output
```

```
##   Feature1 Feature2 Feature3 Feature4
## 1      A      1000      25.5        10
## 2      B      2000      35.5        34
## 3      C      3000      45.5        78
## 4      D      4000      55.5         3
## 5      E      5000      65.5        NA
## 6      F      6000      75.5        NA
## 7      G      7000      85.5        NA
## 8      H      8000      95.5        NA
```

In this case, we can directly use the merge function on both data-frames and set the all argument to TRUE in order to keep all rows from the both data-frames and fill the unavailable values with NA. **only merge feature 1-3. 7/10**

Problem #8

Part A

Import the data from Moodle or shared Google drive, it is called pima.csv. Change the name of the nine columns to preg_times, glucose_test, blood_press, tsk_thickness, serum, bm_index, pedigree_fun, age, class.

```
pima = read.csv("../DATA/pima.csv")
colnames(pima) = c("preg_times",
                  "glucose_test",
                  "blood_press",
                  "tsk_thickness",
                  "serum",
                  "bm_index",
                  "pedigree_fun",
                  "age",
                  "class")
```

This code snippet will simply read the pima dataset from my data folder and change the column names to the requested values.

Part B

All patients (768 Observations) in this dataset contains are females at least 21 years old of Pima Indian heritage. All zero values for the biological variables other than number of times pregnant should be treated as missing values. Count how many zeros are there in each variable (column). For any 0 in the data (except for class and preg_times) assign it as an NA.

```
summary(pima == 0)

##  preg_times      glucose_test      blood_press      tsk_thickness
##  Mode :logical   Mode :logical   Mode :logical   Mode :logical
##  FALSE:657      FALSE:763        FALSE:733        FALSE:541
##  TRUE :111       TRUE :5          TRUE :35         TRUE :227
##    serum        bm_index        pedigree_fun      age
##  Mode :logical   Mode :logical   Mode :logical   Mode :logical
##  FALSE:394       FALSE:757       FALSE:768       FALSE:768
##  TRUE :374       TRUE :11
##    class
##  Mode :logical
##  FALSE:268
##  TRUE :500
```

When we take the summary of the values equal to 0, we will have the dataset converted to TRUE and FALSE logical values. In the summary, we can see how many TRUE values we have which shows us the count of 0 values.

```
exclude.columns = c("class", "preg_times")
selected.cols = !(colnames(pima) %in% exclude.columns)
pima[, selected.cols][pima[, selected.cols] == 0] = NA
```

To replace the 0 values with NA, we should first save the columns we would like to exclude from this process in a vector. We should also create another vector with selected columns, so that we can apply the procedure only to these columns.

After we selected the column we would like to work with, we should use these column names to select the 0 values in these respective columns and replace them with NA value.

```
summary(pima == 0)

##  preg_times      glucose_test      blood_press      tsk_thickness
##  Mode :logical   Mode :logical   Mode :logical   Mode :logical
##  FALSE:657      FALSE:763        FALSE:733        FALSE:541
##  TRUE :111       NA's :5         NA's :35         NA's :227
##    serum        bm_index        pedigree_fun      age
##  Mode :logical   Mode :logical   Mode :logical   Mode :logical
##  FALSE:394       FALSE:757       FALSE:768       FALSE:768
##  NA's :374       NA's :11
```

```
##      class
## Mode :logical
## FALSE:268
##  TRUE :500
```

We can call the summary function again to see that the values are not TRUE anymore and they are set to NA.

Part C

For class variable, check if it is a factor and if not, then make it a factor with levels 0 replaced with neg (for negative diabetic) and 1 replicated with pos (for positive diabetic).

```
class(pima$class)
## [1] "integer"
```

When we check the class of the class variable, we can see that it belongs to the integer class.

```
pima$class = factor(pima$class,
                    levels = c(0, 1),
                    labels = c("neg", "pos"))
```

We can replace this column with a factor using the factor function. In this function, we can determine the levels as 0 and 1. At the same time, we can set the labels of these levels as “neg” and “pos” as requested in the prompt.

```
class(pima$class)
## [1] "factor"
levels(pima$class)
## [1] "neg" "pos"
```

We can check the class and the levels of the variable again and see that it now belongs to the factor class with levels of “pos” and “neg”.

Part D

Make data subsets for four age groups: 21-36, 37-51, 52-66 and 67-81.

```
subset.21.36 = subset(pima, age %in% 21:36)
subset.37.51 = subset(pima, age %in% 37:51)
subset.52.66 = subset(pima, age %in% 52:66)
subset.67.81 = subset(pima, age %in% 67:81)
```

We can directly use the %in% operator to select the age values by checking if they are in a specific range or not. We can also check the answers by looking at the min/max values in the numerical summary of the age column of these subsets.

```
summary(subset.21.36$age)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  21.00   22.00   25.00   26.22   29.00   36.00

summary(subset.37.51$age)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  37.00   39.00   42.00   42.63   45.00   51.00

summary(subset.52.66$age)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  52.00   54.00   58.00   57.91   61.25   66.00

summary(subset.67.81$age)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##       67       67       69       70       70       81
```

Part E

Create a new factor vector called `age.factor`, with `age` in `pima` data replaced with the age group.

```
pima$age.factor[pima$age %in% 21:36] = "21-36"
pima$age.factor[pima$age %in% 37:51] = "37-51"
pima$age.factor[pima$age %in% 52:66] = "52-66"
pima$age.factor[pima$age %in% 67:81] = "67-81"

pima$age.factor = factor(pima$age.factor)
class(pima$age.factor)

## [1] "factor"

summary(pima$age.factor)

## 21-36 37-51 52-66 67-81
##   514   181    64     9
```

In this example we can set the values of the new column `age.factor` to the specific age groups, using the same selection method we used in the previous part.

After setting the values, we can directly convert this column into a factor. Lastly, we can check the class of the `age.factor` to ensure the type is correct. [25/25](#)

Project Problem

Do the following for the approved dataset(s):